student_intervention

October 18, 2016

1 Machine Learning Engineer Nanodegree

1.1 Supervised Learning

1.2 Project 2: Building a Student Intervention System

Welcome to the second project of the Machine Learning Engineer Nanodegree! In this notebook, some template code has already been provided for you, and it will be your job to implement the additional functionality necessary to successfully complete this project. Sections that begin with **'Implementation'** in the header indicate that the following block of code will require additional functionality which you must provide. Instructions will be provided for each section and the specifics of the implementation are marked in the code block with a `'TODO'` statement. Please be sure to read the instructions carefully!

In addition to implementing code, there will be questions that you must answer which relate to the project and your implementation. Each section where you will answer a question is preceded by a **'Question X'** header. Carefully read each question and provide thorough answers in the following text boxes that begin with **'Answer:'**. Your project submission will be evaluated based on your answers to each of the questions and the implementation you provide.

> **Note:** Code and Markdown cells can be executed using the **Shift + Enter** keyboard shortcut. In addition, Markdown cells can be edited by typically double-clicking the cell to enter edit mode.

1.2.1 Question 1 - Classification vs. Regression

Your goal for this project is to identify students who might need early intervention before they fail to graduate. Which type of supervised learning problem is this, classification or regression? Why?

Answer: It's a classification problem , because the label of the data is might need intervention or not ,this is a binary label not a number so it is a classification problem.

1.3 Exploring the Data

Run the code cell below to load necessary Python libraries and load the student data. Note that the last column from this dataset, `'passed'`, will be our target label (whether the student graduated or didn't graduate). All other columns are features about each student.

```
In [17]: # Import libraries
         import numpy as np
         import pandas as pd
```

```
from time import time
from sklearn.metrics import f1_score

# Read student data
student_data = pd.read_csv("student-data.csv")
print "Student data read successfully!"
```

Student data read successfully!

1.3.1 Implementation: Data Exploration

Let's begin by investigating the dataset to determine how many students we have information on, and learn about the graduation rate among these students. In the code cell below, you will need to compute the following: - The total number of students, n_students. - The total number of features for each student, n_features. - The number of those students who passed, n_passed. - The number of those students who failed, n_failed. - The graduation rate of the class, grad_rate, in percent (%).

```
In [18]: # TODO: Calculate number of students
         n_students = student_data.shape[0]

         # TODO: Calculate number of features
         # pay attention to exclude the target label !!
         n_features = student_data.iloc[:,:-1].shape[1]

         # TODO: Calculate passing students
         n_passed = n_passed = student_data[student_data.passed == 'yes'].shape[0]

         # TODO: Calculate failing students
         n_failed = n_students - n_passed

         # TODO: Calculate graduation rate
         grad_rate = float(n_passed)/n_students*100

         # Print the results
         print "Total number of students: {}".format(n_students)
         print "Number of features: {}".format(n_features)
         print "Number of students who passed: {}".format(n_passed)
         print "Number of students who failed: {}".format(n_failed)
         print "Graduation rate of the class: {:.2f}%".format(grad_rate)
```

```
Total number of students: 395
Number of features: 30
Number of students who passed: 265
Number of students who failed: 130
Graduation rate of the class: 67.09%
```

2

1.4 Preparing the Data

In this section, we will prepare the data for modeling, training and testing.

1.4.1 Identify feature and target columns

It is often the case that the data you obtain contains non-numeric features. This can be a problem, as most machine learning algorithms expect numeric data to perform computations with.

Run the code cell below to separate the student data into feature and target columns to see if any features are non-numeric.

```
In [19]: # Extract feature columns
         feature_cols = list(student_data.columns[:-1])

         # Extract target column 'passed'
         target_col = student_data.columns[-1]

         # Show the list of columns
         print "Feature columns:\n{}".format(feature_cols)
         print "\nTarget column: {}".format(target_col)

         # Separate the data into feature data and target data (X_all and y_all, re
         X_all = student_data[feature_cols]
         y_all = student_data[target_col]

         # Show the feature information by printing the first five rows
         print "\nFeature values:"
         print X_all.head()

Feature columns:
['school', 'sex', 'age', 'address', 'famsize', 'Pstatus', 'Medu', 'Fedu', 'Mjob', '

Target column: passed

Feature values:
  school sex  age address famsize Pstatus  Medu  Fedu      Mjob      Fjob  \
0     GP   F   18       U     GT3       A     4     4   at_home   teacher
1     GP   F   17       U     GT3       T     1     1   at_home     other
2     GP   F   15       U     LE3       T     1     1   at_home     other
3     GP   F   15       U     GT3       T     4     2    health  services
4     GP   F   16       U     GT3       T     3     3     other     other

      ...   higher internet  romantic  famrel  freetime goout Dalc Walc health  \
0     ...      yes       no        no       4         3     4    1    1      3
1     ...      yes      yes        no       5         3     3    1    1      3
2     ...      yes      yes        no       4         3     2    2    3      3
3     ...      yes      yes       yes       3         2     2    1    1      5
4     ...      yes       no        no       4         3     2    1    2      5
```

```
   absences
0         6
1         4
2        10
3         2
4         4

[5 rows x 30 columns]
```

1.4.2   Preprocess Feature Columns

As you can see, there are several non-numeric columns that need to be converted! Many of them are simply yes/no, e.g. internet. These can be reasonably converted into 1/0 (binary) values.

Other columns, like Mjob and Fjob, have more than two values, and are known as *categorical variables*. The recommended way to handle such a column is to create as many columns as possible values (e.g. Fjob_teacher, Fjob_other, Fjob_services, etc.), and assign a 1 to one of them and 0 to all others.

These generated columns are sometimes called *dummy variables*, and we will use the pandas.get_dummies() function to perform this transformation. Run the code cell below to perform the preprocessing routine discussed in this section.

```
In [20]: def preprocess_features(X):
             ''' Preprocesses the student data and converts non-numeric binary var
                 binary (0/1) variables. Converts categorical variables into dummy

             # Initialize new output DataFrame
             output = pd.DataFrame(index = X.index)

             # Investigate each feature column for the data
             for col, col_data in X.iteritems():

                 # If data type is non-numeric, replace all yes/no values with 1/0
                 if col_data.dtype == object:
                     col_data = col_data.replace(['yes', 'no'], [1, 0])

                 # If data type is categorical, convert to dummy variables
                 if col_data.dtype == object:
                     # Example: 'school' => 'school_GP' and 'school_MS'
                     col_data = pd.get_dummies(col_data, prefix = col)

                 # Collect the revised columns
                 output = output.join(col_data)

             return output

         X_all = preprocess_features(X_all)
         print "Processed feature columns ({} total features):\n{}".format(len(X_al
```

4

```
Processed feature columns (48 total features):
['school_GP', 'school_MS', 'sex_F', 'sex_M', 'age', 'address_R', 'address_U', 'fams
```

1.4.3   Implementation: Training and Testing Data Split

So far, we have converted all *categorical* features into numeric values. For the next step, we split the data (both features and corresponding labels) into training and test sets. In the following code cell below, you will need to implement the following: - Randomly shuffle and split the data (X_all, y_all) into training and testing subsets. - Use 300 training points (approximately 75%) and 95 testing points (approximately 25%). - Set a random_state for the function(s) you use, if provided. - Store the results in X_train, X_test, y_train, and y_test.

```python
In [25]:  # TODO: Import any additional functionality you may need here
          from sklearn.model_selection import train_test_split

          # TODO: Set the number of training points
          num_train = 300

          # Set the number of testing points
          num_test = X_all.shape[0] - num_train

          # TODO: Shuffle and split the dataset into the number of training and test
          # use stratify attribute to ensure the ratio between the two classes is th
          X_train, X_test, y_train, y_test = train_test_split(X_all, y_all, test_siz

          # Show the results of the split
          print "Training set has {} samples.".format(X_train.shape[0])
          print "Testing set has {} samples.".format(X_test.shape[0])

Training set has 300 samples.
Testing set has 95 samples.
```

1.5   Training and Evaluating Models

In this section, you will choose 3 supervised learning models that are appropriate for this problem and available in scikit-learn. You will first discuss the reasoning behind choosing these three models by considering what you know about the data and each model's strengths and weaknesses. You will then fit the model to varying sizes of training data (100 data points, 200 data points, and 300 data points) and measure the F1 score. You will need to produce three tables (one for each model) that shows the training set size, training time, prediction time, F1 score on the training set, and F1 score on the testing set.

The following supervised learning models are currently available in scikit-learn **that you may choose from:** - Gaussian Naive Bayes (GaussianNB) - Decision Trees - Ensemble Methods (Bagging, AdaBoost, Random Forest, Gradient Boosting) - K-Nearest Neighbors (KNeighbors) - Stochastic Gradient Descent (SGDC) - Support Vector Machines (SVM) - Logistic Regression

1.5.1 Question 2 - Model Application

List three supervised learning models that are appropriate for this problem. For each model chosen - Describe one real-world application in industry where the model can be applied. *(You may need to do a small bit of research for this — give references!)* - What are the strengths of the model; when does it perform well? - What are the weaknesses of the model; when does it perform poorly? - What makes this model a good candidate for the problem, given what you know about the data?

Answer: Decision Trees can be used in oil field,see reference as http://ardent.mit.edu/real_options/Real_opts_papers/Babajide_Thesis_FINAL.pdf The advantage of the decision trees is they are simple to understand and interpret,you can also have value even with little hard data.Calculations can get very complex particularly if many values are uncertain is their weakness.Since their isn't a large amount of data in this project,so I think a light-weighted machine learning model would be better,therefore I choose decision trees as well as ensemble methods and logistic regression.In a real industrial scenario,classification of bad accounts in credit card uses the gradient boosting model see as http://cs229.stanford.edu/proj2014/Chengwei%20Yuan,%20Classification%20of%20Bad%20Accounts%20in%2 The strength is that even you run the model over time,testing error is still decrease, according to the udacity machine learning video.The weakness is it is more complex than decision tree and the parameters may not be tuned easily.Why I choose this model as a candidate is explained before.Logistic regression can be used in retail industry,http://xueshu.baidu.com/s?wd=paperuri:(69c3fd9dbc2e406d897abd473c0cbac6)&filter=sc_long_sign& 8&sc_us=7649060419404245118 shows this.They can be easily interpreted and fast to train but need to transform the features if they are not linear seperable.

1.5.2 Setup

Run the code cell below to initialize three helper functions which you can use for training and testing the three supervised learning models you've chosen above. The functions are as follows: - `train_classifier` - takes as input a classifier and training data and fits the classifier to the data. - `predict_labels` - takes as input a fit classifier, features, and a target labeling and makes predictions using the F1 score. - `train_predict` - takes as input a classifier, and the training and testing data, and performs `train_clasifier` and `predict_labels`. - This function will report the F1 score for both the training and testing data separately.

```python
In [22]: def train_classifier(clf, X_train, y_train):
             ''' Fits a classifier to the training data. '''

             # Start the clock, train the classifier, then stop the clock
             start = time()
             clf.fit(X_train, y_train)
             end = time()

             # Print the results
             print "Trained model in {:.4f} seconds".format(end - start)


         def predict_labels(clf, features, target):
             ''' Makes predictions using a fit classifier based on F1 score. '''
```

```
            # Start the clock, make predictions, then stop the clock
            start = time()
            y_pred = clf.predict(features)
            end = time()

            # Print and return results
            print "Made predictions in {:.4f} seconds.".format(end - start)
            return f1_score(target.values, y_pred, pos_label='yes')


        def train_predict(clf, X_train, y_train, X_test, y_test):
            ''' Train and predict using a classifer based on F1 score. '''

            # Indicate the classifier and the training set size
            print "Training a {} using a training set size of {}. . .".format(clf.

            # Train the classifier
            train_classifier(clf, X_train, y_train)

            # Print the results of prediction for both training and testing
            print "F1 score for training set: {:.4f}.".format(predict_labels(clf,
            print "F1 score for test set: {:.4f}.".format(predict_labels(clf, X_te
```

1.5.3   Implementation: Model Performance Metrics

With the predefined functions above, you will now import the three supervised learning models of your choice and run the `train_predict` function for each one. Remember that you will need to train and predict on each classifier for three different training set sizes: 100, 200, and 300. Hence, you should expect to have 9 different outputs below — 3 for each model using the varying training set sizes. In the following code cell, you will need to implement the following: - Import the three supervised learning models you've discussed in the previous section. - Initialize the three models and store them in `clf_A`, `clf_B`, and `clf_C`. - Use a `random_state` for each model you use, if provided. - **Note:** Use the default settings for each model — you will tune one specific model in a later section. - Create the different training set sizes to be used to train each model. - *Do not reshuffle and resplit the data! The new training points should be drawn from* `X_train` *and* `y_train`. - Fit each model with each training set size and make predictions on the test set (9 in total).
Note: Three tables are provided after the following code cell which can be used to store your results.

```
In [23]: # TODO: Import the three supervised learning models from sklearn
         from sklearn.tree import DecisionTreeClassifier
         from sklearn.ensemble import GradientBoostingClassifier
         from sklearn.linear_model import LogisticRegression

         # TODO: Initialize the three models
         clf_A = DecisionTreeClassifier(random_state=0)
         clf_B = GradientBoostingClassifier(random_state=0)
         clf_C = LogisticRegression(random_state=0)
```

7

```python
        # TODO: Set up the training set sizes
        X_train_100 = X_train[:100]
        y_train_100 = y_train[:100]

        X_train_200 = X_train[:200]
        y_train_200 = y_train[:200]

        X_train_300 = X_train
        y_train_300 = y_train

        # TODO: Execute the 'train_predict' function for each classifier and each
        # train_predict(clf_A, X_train_100, y_train_100, X_test, y_test)
        # train_predict(clf_A, X_train_200, y_train_200, X_test, y_test)
        # train_predict(clf_A, X_train_300, y_train_300, X_test, y_test)

        # train_predict(clf_B, X_train_100, y_train_100, X_test, y_test)
        # train_predict(clf_B, X_train_200, y_train_200, X_test, y_test)
        # train_predict(clf_B, X_train_300, y_train_300, X_test, y_test)

        # train_predict(clf_C, X_train_100, y_train_100, X_test, y_test)
        # train_predict(clf_C, X_train_200, y_train_200, X_test, y_test)
        # train_predict(clf_C, X_train_300, y_train_300, X_test, y_test)

        # use for loop to simply the code
        for clf in [clf_A,clf_B,clf_C]:
            for size in [100,200,300]:
                train_predict(clf,X_train[:size],y_train[:size],X_test,y_test)
```

```
Training a DecisionTreeClassifier using a training set size of 100. . .
Trained model in 0.0030 seconds
Made predictions in 0.0010 seconds.
F1 score for training set: 1.0000.
Made predictions in 0.0010 seconds.
F1 score for test set: 0.7107.
Training a DecisionTreeClassifier using a training set size of 200. . .
Trained model in 0.0020 seconds
Made predictions in 0.0000 seconds.
F1 score for training set: 1.0000.
Made predictions in 0.0000 seconds.
F1 score for test set: 0.6885.
Training a DecisionTreeClassifier using a training set size of 300. . .
Trained model in 0.0020 seconds
Made predictions in 0.0000 seconds.
F1 score for training set: 1.0000.
Made predictions in 0.0000 seconds.
F1 score for test set: 0.6875.
Training a GradientBoostingClassifier using a training set size of 100. . .
```

```
Trained model in 0.0480 seconds
Made predictions in 0.0010 seconds.
F1 score for training set: 1.0000.
Made predictions in 0.0010 seconds.
F1 score for test set: 0.7360.
Training a GradientBoostingClassifier using a training set size of 200. . .
Trained model in 0.0640 seconds
Made predictions in 0.0010 seconds.
F1 score for training set: 1.0000.
Made predictions in 0.0010 seconds.
F1 score for test set: 0.7794.
Training a GradientBoostingClassifier using a training set size of 300. . .
Trained model in 0.0870 seconds
Made predictions in 0.0020 seconds.
F1 score for training set: 0.9710.
Made predictions in 0.0000 seconds.
F1 score for test set: 0.7714.
Training a LogisticRegression using a training set size of 100. . .
Trained model in 0.0280 seconds
Made predictions in 0.0180 seconds.
F1 score for training set: 0.8759.
Made predictions in 0.0010 seconds.
F1 score for test set: 0.7500.
Training a LogisticRegression using a training set size of 200. . .
Trained model in 0.0030 seconds
Made predictions in 0.0000 seconds.
F1 score for training set: 0.8532.
Made predictions in 0.0000 seconds.
F1 score for test set: 0.8085.
Training a LogisticRegression using a training set size of 300. . .
Trained model in 0.0040 seconds
Made predictions in 0.0000 seconds.
F1 score for training set: 0.8402.
Made predictions in 0.0010 seconds.
F1 score for test set: 0.7770.
```

1.5.4 Tabular Results

Edit the cell below to see how a table can be designed in Markdown. You can record your results from above in the tables provided.

** Classifer 1 - Decision Tree**

Training Set Size	Training Time	Prediction Time (test)	F1 Score (train)	F1 Score (test)
100	0.003s	0.001s	1.0	0.7107
200	0.002s	0.000s	1.0	0.6885
300	0.002s	0.000s	1.0	0.6875

Training Set Size	Training Time	Prediction Time (test)	F1 Score (train)	F1 Score (test)

** Classifer 2 - Gradient Boosting**

Training Set Size	Training Time	Prediction Time (test)	F1 Score (train)	F1 Score (test)
100	0.0480s	0.001s	1.0	0.7360
200	0.0640s	0.001s	1.0	0.7794
300	0.0870s	0.0s	0.9710	0.7714

** Classifer 3 - Logistic Regression**

Training Set Size	Training Time	Prediction Time (test)	F1 Score (train)	F1 Score (test)
100	0.0280s	0.001s	0.8759	0.7500
200	0.003s	0.0s	0.8532	0.8085
300	0.004s	0.001s	0.8402	0.7770

1.6 Choosing the Best Model

In this final section, you will choose from the three supervised learning models the *best* model to use on the student data. You will then perform a grid search optimization for the model over the entire training set (`X_train` and `y_train`) by tuning at least one parameter to improve upon the untuned model's F1 score.

1.6.1 Question 3 - Choosing the Best Model

Based on the experiments you performed earlier, in one to two paragraphs, explain to the board of supervisors what single model you chose as the best model. Which model is generally the most appropriate based on the available data, limited resources, cost, and performance?

Answer: Based on the experiments I performed earlier,I choose the logistic regression model as the best model.When the training set size is 300,its F1 score is highest among all three models and it trains very fast,so that I can deploy it quickly to see results and make it a good benchmark if I want to discover some other models to outperform it.

1.6.2 Question 4 - Model in Layman's Terms

In one to two paragraphs, explain to the board of directors in layman's terms how the final model chosen is supposed to work. Be sure that you are describing the major qualities of the model, such as how the model is trained and how the model makes a prediction. Avoid using advanced mathematical or technical jargon, such as describing equations or discussing the algorithm implementation.

Answer: Logistic Regression is a machine learning model which uses logistic function to map the input features to ouput target values.Although there are some maths details in it , you can easily use it from sklearn.Import the logistic regression model ,use training dataset to fit

the model , after that params in that model are fixed.Then you can use predict method to predict the result for the data you've never seen.

1.6.3   Implementation: Model Tuning

Fine tune the chosen model. Use grid search (GridSearchCV) with at least one important parameter tuned with at least 3 different values. You will need to use the entire training set for this. In the code cell below, you will need to implement the following: - Import sklearn.grid_search.gridSearchCV and sklearn.metrics.make_scorer. - Create a dictionary of parameters you wish to tune for the chosen model. - Example: parameters = {'parameter' : [list of values]}. - Initialize the classifier you've chosen and store it in clf. - Create the F1 scoring function using make_scorer and store it in f1_scorer. - Set the pos_label parameter to the correct value! - Perform grid search on the classifier clf using f1_scorer as the scoring method, and store it in grid_obj. - Fit the grid search object to the training data (X_train, y_train), and store it in grid_obj.

```
In [24]:  # TODO: Import 'GridSearchCV' and 'make_scorer'
          import numpy
          from sklearn.grid_search import GridSearchCV
          from sklearn.metrics import make_scorer

          # TODO: Create the parameters list you wish to tune
          Cs = numpy.array([0.01,0.1,1.0,10.0])
          parameters = dict(C=Cs)

          # TODO: Initialize the classifier
          clf = LogisticRegression(random_state=1)

          # TODO: Make an f1 scoring function using 'make_scorer'
          f1_scorer = make_scorer(f1_score,pos_label='yes')

          # TODO: Perform grid search on the classifier using the f1_scorer as the s
          grid_obj = GridSearchCV(estimator=clf,param_grid=parameters,scoring=f1_sco

          # TODO: Fit the grid search object to the training data and find the optim
          grid_obj = grid_obj.fit(X_train,y_train)

          # Get the estimator
          clf = grid_obj.best_estimator_

          # Report the final F1 score for training and testing after parameter tunin
          print "Tuned model has a training F1 score of {:.4f}.".format(predict_labe
          print "Tuned model has a testing F1 score of {:.4f}.".format(predict_label

Made predictions in 0.0000 seconds.
Tuned model has a training F1 score of 0.8157.
Made predictions in 0.0000 seconds.
Tuned model has a testing F1 score of 0.8129.
```

1.6.4   Question 5 - Final F1 Score

What is the final model's F1 score for training and testing? How does that score compare to the untuned model?

   Answer: traing f1 score is 0.8157,while test f1 score is 0.8129.Its the highest test f1 score compared to the untuned Logistic Regression model.Since I forgot to use the stratify attribute of train_test_split,I got worse result last time.But now , the result seems to be reasonable.

> **Note**: Once you have completed all of the code implementations and successfully answered each question above, you may finalize your work by exporting the iPython Notebook as an HTML document. You can do this by using the menu above and navigating to
> **File -> Download as -> HTML (.html)**. Include the finished document along with this notebook as your submission.